


KNOWLEDGE WORKER EDITION

Daily AI Brief Command Center

A practical guide to monitoring publication, understanding editorial evidence, reviewing quality, and keeping every high-impact decision human-controlled.



The Command Center turns GitHub evidence into an understandable, read-only operational view.

AUDIENCE	GUIDE	SITE
Knowledge workers using AI	Version 1.0 • Sep 2026	Owner-only Command Center

[Open the Daily AI Brief Command Center](#)

START HERE

How to use this guide

Read Chapters 1-3 first. Then jump directly to the view or decision you need.

IF YOU NEED TO...	GO TO	YOU WILL LEARN
Understand the system	Chapters 1-2	What the Command Center controls, where evidence comes from, and what remains public.
Run a morning check	Chapters 3-5	How to scan Overview and trace Today's Run.
Judge editorial quality	Chapters 6-9	Selection, reader evidence, QA, incidents, and corrections.
Explore the system	Chapters 10-12	History, components, automations, and Trend Radar.
Make a safe decision	Chapters 13-15	Evidence language, approval boundaries, and operating rhythm.
Solve a problem	Appendices	Troubleshooting, glossary, and the one-page checklist.

LEARNING PATTERN Each chapter follows concept → interpretation → action → what good looks like. You do not need to understand code, Git, or automation design.

Five rules worth remembering

1. **GitHub is authoritative** - The dashboard explains canonical records; it does not replace them.
2. **Missing is not healthy** - Unavailable evidence stays Pending, Degraded, or Not measured.
3. **Green is evidence-bounded** - A passing card means the configured checks passed - not that every imaginable risk is gone.
4. **History is append-only** - Corrections preserve prior observations and add the resolution.
5. **Silence is not approval** - High-impact production and editorial changes require explicit owner intent.

01

What the Command Center does

See the Daily AI Brief as a governed publishing system - not just a webpage.

The Daily AI Brief Command Center is a private, owner-only operational console. It brings today’s edition, editorial selection, quality evidence, reader signals, incidents, history, automations, and trends into one read-only view. The public brief, story pages, archive, feeds, feedback, QA history, and Trend Radar remain separate public assets.


LAYER	SOURCE OF TRUTH	COMMAND CENTER JOB
Editorial	Canonical dated edition and candidate records	Explain why six stories were selected.
Quality	QA, accessibility, CI, and deployment records	Surface pass, failure, repair, and evidence links.
Readers	Privacy-safe aggregate analytics and ratings	Show signal strength without overstating a small audience.
Operations	Automation summaries and append-only ledgers	Expose status and history without changing configuration.
Governance	Program state, policy profile, and owner decisions	Keep high-impact actions outside the dashboard.

PLAIN LANGUAGE Use the Command Center to answer five questions: Did today publish? Are the six stories valid? Did QA pass? Is anything broken? Does anything need my decision?

02

The source-of-truth model

Understand why “Live GitHub records” matters before reading any score.



Live Overview: evidence status, edition, selection, feedback, issues, and owner decision.

The evidence indicator shows whether the page loaded current repository records or a timestamped verified fallback. If live retrieval fails, a safe fallback may appear, but the overall state cannot be presented as fully healthy.

SOURCE LABEL	MEANING	ACTION
Live GitHub records	Current canonical records loaded successfully.	Use timestamps and evidence links for normal review.
Verified fallback	A known snapshot is shown because live retrieval failed.	Treat as degraded; verify before any decision.
Data unavailable	Neither live records nor a usable fallback could load.	Do not infer publication or quality state.

03

Your 10-minute quick start

Build confidence with one safe scan from publication to evidence.

Minute 0-2: orient

1. Confirm the evidence label and refresh time.
2. Read Current edition, Editorial selection, Operational issues, and Owner decision separately.
3. Open Needs attention before following a green total.

Minute 2-5: trace the edition

- Open Today.
- Confirm 20 candidates narrowed to six selected stories with exact 2/2/2 focus balance.
- Confirm deterministic CI, GitHub Pages, independent QA, and accessibility evidence.

Minute 5-8: inspect quality and corrections

- Open Quality and read the initial and final QA states.
- Open Incidents for current findings and the approval queue.
- Open Corrections to distinguish unique corrections from immutable ledger rows.

Minute 8-10: close the loop

- If no action is needed, leave the system unchanged.
- If evidence is stale or missing, request a targeted verification.
- If a high-impact change is proposed, decide in ChatGPT Work - not in the dashboard.

WHAT GOOD LOOKS LIKE You can explain what published, why the stories were chosen, what evidence passed, what remains uncertain, and whether a human decision is required.

04

Read the Overview

Turn five summary cards into a reliable priority order.

CARD	READ IT AS	DO NOT ASSUME
Current edition	Freshness of the latest canonical edition.	That the public page is correct without deployment evidence.
Editorial selection	Selected count and 2/2/2 balance.	That popularity can override evidence or novelty gates.
Reader feedback	Progress toward the learning evidence threshold.	That zero ratings is a defect.
Operational issues	Current findings derived from available records.	That “None” proves every subsystem is healthy.
Owner decision	Whether a bounded recommendation needs review.	That the dashboard can execute the decision.


Reliable triage order

1. **Evidence source** - Fallback or unavailable data changes how every other card should be read.
2. **Owner decision** - A prepared recommendation may be waiting.
3. **Operational issues** - Failed or degraded evidence can affect today’s edition.
4. **Current edition** - Freshness determines whether you are reviewing today or yesterday.
5. **Editorial and reader signals** - Interpret them after hard operational gates.

05

Trace Today's Run

Follow the current edition from evidence-screened candidates through publication and QA.



The editorial funnel preserves evidence quality and exact focus-area balance.

CHECKPOINT	WHAT TO VERIFY	EVIDENCE LINK
Candidate selection	Candidate count, selected count, scoring rationale, policy profile.	Candidate record
Deterministic CI	Latest check conclusion matches the observed repository head.	Workflow evidence
GitHub Pages	Deployment completed for the intended revision.	Deployment evidence
Independent QA	Initial result, repairs, final result, and defects.	QA record
Accessibility	Story alternatives and disposition were reviewed.	Accessibility record
Public edition	Permanent story pages and dated brief open correctly.	Current brief

KEY DISTINCTION A successful deployment proves a revision reached GitHub Pages. It does not, by itself, prove the editorial record or rendered content is correct.

06

Understand Editorial Intelligence

See why the six stories earned a place without turning scoring into a black box.

Editorial Intelligence summarizes the edition’s evidence mix, novelty disposition, focus allocation, and selection record. It links back to the canonical edition and candidate rationale so that every summary can be checked.

SIGNAL	WHY IT EXISTS	HARD BOUNDARY
Evidence class	Distinguishes official announcements, research, practitioner analysis, and other support.	Weak evidence cannot be rescued by popularity.
Novelty disposition	Protects the 30-day novelty gate and makes repeats visible.	Repetition requires an explicit supported exception.
2/2/2 allocation	Balances Technical AI Engineering, Applied GenAI, and Agents for Non-Technical People.	The exact allocation remains protected.
Candidate rationale	Explains why selected stories outranked alternatives.	Scores do not override accessibility or privacy.


INTERPRETATION Editorial scoring supports judgment. Evidence, novelty, accessibility, privacy, and exact focus balance are hard gates.

07

Use Readers & Editorial Learning

Learn from a small audience without pretending weak signals are strong evidence.

Editorial learning: evidence before activation



Every learning threshold must pass before a separate owner approval.

Reader analytics are directional: small counts are suppressed, shares count initiations, and ratings stay anonymous. Ratings may recommend a small Practical Value adjustment; they cannot weaken evidence, novelty, accessibility, privacy, or focus balance. Learning inactivity is expected until every threshold is met; it is not a defect.

THRESHOLD	MINIMUM	WHY
Fully rated editions	5	Prevents one unusual day from dominating.
Total story ratings	30	Creates a minimally useful signal base.
Ratings per focus	10 each	Protects balanced learning across all three areas.
Candidate-pool comparisons	5 passing sets	Checks that the proposed weighting improves selection quality.
Owner decision	Explicit approval	Keeps editorial authority human-controlled.

08

Interpret Quality & QA

Separate initial defects, automated repairs, final validation, and deployment proof.

FIELD	QUESTION
Initial QA	Did the first validation pass before any repair?
Repairs	What was changed, and was each change bounded to the defect?
Final QA	Did the complete required check set pass after repair?
Accessibility	Were alternatives, structure, links, and readable presentation checked?
Deployment evidence	Did the validated revision reach the public site?
Public QA history	Can a reader inspect the edition's quality record?

FAILURE RULE A critical publication failure forces Failed even when every other available indicator is green.


When QA repairs something

1. Read the initial failure and severity.
2. Verify the repair is specific to the cause.
3. Confirm final QA reran the full relevant checks.
4. Confirm the deployed revision matches the validated revision.
5. Keep the repair and result in append-only history.

09

Handle incidents and corrections

Move from a signal to a decision-ready assessment without erasing history.



A controlled path turns a detected issue into evidence, validation, a human decision, and an auditable outcome.

VIEW	USE IT FOR	IMPORTANT DISTINCTION
Incidents	Current findings, impact, evidence, proposed action, validation, rollback, and approval state.	No ledger entries does not prove health.
Corrections	Material factual changes and affected public pages.	Unique corrections differ from immutable ledger rows.
Approval queue	Decision-ready high-impact changes.	No action is enabled inside the Command Center.

APPEND-ONLY RULE When duplicate observations or a corrected fact exist, preserve the original rows and append reconciliation evidence. Report canonical event counts without deleting history.

10

Search Story & Edition History

Use local search to explore the archive without sending your query to another service.

The History view lets you search by story, company, or topic and filter by focus area. Results show date, permanent story page, focus, evidence class, and availability. The current implementation can display up to 125 indexed stories.

TASK	ACTION	WHAT GOOD LOOKS LIKE
Find prior coverage	Search a company, product, or concept.	You see dated permanent pages and evidence types.
Check focus balance over time	Filter one focus area.	The list narrows without losing dates or evidence labels.
Investigate repetition	Compare related headlines and dates.	Novelty and prior-coverage evidence can be traced.
Share a story	Open the permanent page.	The URL remains stable beyond the daily homepage.


PRIVACY Search filtering happens in your browser. The query text is not sent to a search service.

11

Read Components and Automations

Understand subsystem health and schedules without exposing private configuration.

The morning operating rhythm



All times America/Chicago. The verification step cannot publish, repair, approve, or change production.

The publisher, QA, and read-only verification run in sequence each morning.

AREA	WHAT YOU CAN SEE	WHAT YOU CANNOT DO
Components	Publisher, ranker, Story Memory, CI, Pages, QA, analytics, learning, and feeds.	Change a subsystem from the dashboard.
Automations	Name, state, schedule, last verified run, and safe scope summary.	Pause, resume, edit, or reveal private instructions.
Evidence links	Public records and workflow results.	Access credentials, tokens, raw headers, or private prompts.

OPERATING BOUNDARY The 7:40 AM Command Center validation is intentionally unable to repair, publish, approve, or change production.

12

Use Trend Radar responsibly

Treat trend classification as descriptive evidence - not a forecast or editorial shortcut.

FIELD	MEANING	SAFE INTERPRETATION
Classification	Descriptive movement across recent evidence windows.	A pattern in observed stories, not a prediction.
Confidence	Strength and consistency of the supporting record.	Confidence in classification, not certainty about the future.
Evidence count	Supporting stories across the configured windows.	Volume, not importance by itself.
Linked record	Dated trend evidence in GitHub.	The place to verify the summary.

EDITORIAL SAFEGUARD Trend prominence can guide attention, but it does not override story-level evidence, novelty, or focus allocation.

13

Read evidence states correctly

Use explicit language so missing data never turns into false reassurance.

STATE	MEANING	NEXT MOVE
Healthy / Pass	Configured current evidence passed.	Continue normal monitoring.
Pending	Evidence is expected but incomplete or not yet due.	Wait or request the smallest needed check.
Degraded / At risk	A limitation, stale input, or failed noncritical signal exists.	Open the evidence and assess impact.
Failed	A critical publication or hard-gate condition failed.	Stop assumptions; investigate and recover.
Not measured	No trustworthy metric exists for this field.	Do not substitute zero.
Unspecified	Legacy records lack the modern field.	Keep the limitation visible.

GOLDEN RULE The Command Center must never create a green card from missing evidence.

14

Keep high-impact decisions human

Know which actions stay in ChatGPT Work and how to give an unambiguous answer.

DECISION-READY FIELD	WHAT TO EXPECT
Error or proposal	Exact behavior, record, route, or policy change.
Assessment	Impact, affected audience, cause, and confidence.
Corrective action	The smallest prepared change that addresses the cause.
Validation	QA, link, accessibility, publication, and regression evidence.
Recovery	Rollback or tested correction if the change regresses.
Decision	A direct request to approve, hold, or reject the specific scope.

Clear decision language

- **Approve production** - Authorize only the named version and stated scope.
- **Do not publish** - Reject the proposed production move.
- **Hold - verify X** - Name the one missing fact that must be checked before deciding.

HUMAN CONTROL A notification, acknowledgement, question, or silence is not authorization.

15

Adopt a sustainable operating rhythm

Keep oversight light enough to use every day and rigorous enough to trust.

CADENCE	REVIEW	SUCCESS SIGNAL
Daily after 7:40 AM	Source label, edition, QA, issues, owner decision.	Today’s state is explained in under ten minutes.
Weekly	Reader evidence, learning thresholds, corrections, automation freshness.	Every stale or unusual signal has a follow-up.
Monthly	Trend evidence, archive patterns, recurring defects, policy fitness.	Recommendations remain bounded and evidence-backed.
After an incident	Cause, impact, correction, validation, decision, outcome.	History tells the complete story end to end.
Before any high-impact change	Scope, evidence, rollback, protected dimensions.	Explicit approve, hold, or reject language exists.

16

Troubleshoot common symptoms

Choose the smallest safe next check before changing anything.

SYMPTOM	LIKELY INTERPRETATION	NEXT CHECK
Overview still shows yesterday	Today's canonical record may be missing, stale, or before the expected window.	Refresh, check source label, then open Today evidence.
Live GitHub changed to fallback	Repository retrieval failed; a verified snapshot is displayed.	Treat as degraded and recheck GitHub evidence.
Six stories but QA is Pending	Editorial selection exists; quality evidence is incomplete.	Open QA and deployment records before trusting publication.
No operational issues	No finding was derived from current records.	Do not infer that every subsystem was audited.
Zero reader ratings	Learning evidence is insufficient.	Continue rating; do not activate weighting.
Immutable rows exceed unique corrections	Duplicate or reconciliation rows are preserved.	Use canonical unique count and read reconciliation evidence.
Automation appears stale	Safe snapshot or last-verified time may be old.	Verify the authoritative schedule outside the dashboard.
A public page looks wrong	Deployment success and content correctness differ.	Open permanent page, edition record, QA, and correction history.

17

Glossary

Use shared language to keep operational decisions precise.

TERM	DEFINITION
Canonical edition record	The authoritative dated machine record for the selected stories and edition facts.
Candidate pool	The evidence-screened set considered before the six-story selection.
Deterministic CI	Repeatable automated checks applied to the observed repository revision.
Evidence link	A direct path to the supporting public record or workflow result.
Hard gate	A requirement that scoring or popularity cannot override.
Owner decision	Explicit approval, hold, or rejection made outside the read-only dashboard.
Permanent story page	A stable public URL for one story after it leaves the homepage.
Policy profile	The named publication and editorial rule set used for the edition.
Source mode	Whether evidence came from live GitHub records or a verified fallback.
Story Memory	The record that supports novelty and prior-coverage checks.
Trend Radar	Descriptive classification of observed story patterns across evidence windows.

18

One-page operator checklist

Print this page or keep it beside the Command Center.

SCAN

- ☐ Confirm source label and refresh time
- ☐ Confirm current edition and six-story selection
- ☐ Review operational issues and owner decision

VERIFY

- ☐ Trace candidate, CI, deployment, QA, and accessibility evidence
- ☐ Read Pending, Degraded, Failed, Not measured, and Unspecified literally
- ☐ Check corrections and incident history before assuming closure

INTERPRET

- ☐ Separate deployment success from content correctness
- ☐ Treat reader analytics as directional
- ☐ Keep hard editorial gates separate from popularity

DECIDE

- ☐ Review scope, impact, validation, and recovery
- ☐ Use Approve production / Do not publish / Hold - verify X
- ☐ Confirm the outcome is appended to history

OPERATOR PROMISE No guessed facts. No silent promotion. No erased history. No high-impact action without explicit owner intent.

19

Sources and document control

Know what is universal, what is implementation-specific, and when to review this guide.

Primary implementation sources

[Daily AI Brief repository](#) Canonical publication code, records, ledgers, public routes, and documentation.

[Public Daily AI Brief](#) Current edition, archive, story pages, QA, Trend Radar, feedback, and feeds.

[Daily AI Brief Command Center](#) Private owner-only operational console documented in this guide.

This guide describes the observed Command Center implementation and its evidence contracts as of September 2026. Counts, dates, and current states will change. The reading method, source hierarchy, privacy boundary, append-only history, and approval gate remain the durable operating principles.

DOCUMENT RECORD Daily AI Brief Command Center User Guide • Knowledge Worker Edition • Version 1.0 • Prepared September 2026 • Review after a major navigation, evidence-contract, approval, privacy, or automation change.